



SAMSUNG


Java Concurrency and Multithreading

TS. Nguyễn Văn Sơn
BM Công Nghệ Phần Mềm – FIT UET

Bài giảng được tài trợ bởi
Trung tâm Nghiên cứu và Phát triển Samsung Việt Nam (SRV)

Thông tin giảng viên

- Liên lạc: Phòng 321, E3, sonnguyen@vnu.edu.vn
- Nghiên cứu: AI for Software Engineering, Automated AI Engineering, Quality Assurance AI-enabled Systems
- Một số bài toán điển hình:
 - Tự động gợi ý/sinh mã nguồn
 - Phát hiện lỗ hổng bảo mật phần mềm
 - Tự động hóa quá trình phát triển giải pháp dựa trên AI
 - Đảm bảo chất lượng hệ thống xe tự hành
- Nhóm NC: Intelligent Software Engineering (iSE)



Tôi, năm 2022

Contents


- Concurrency and Multithreading
- Thread and Runnable in Java
- Executor Service & Thread Pool
- Callable and Future
- Thread Synchronization
- Locks and Atomic Variables
- Virtual Thread

Core concepts

Concurrency

- **Concurrency**: run several programs or several parts of a program in parallel
- A task can be performed asynchronously or in parallel
 - Improves the throughput and the interactivity of the program.

Concurrency vs. Parallelism

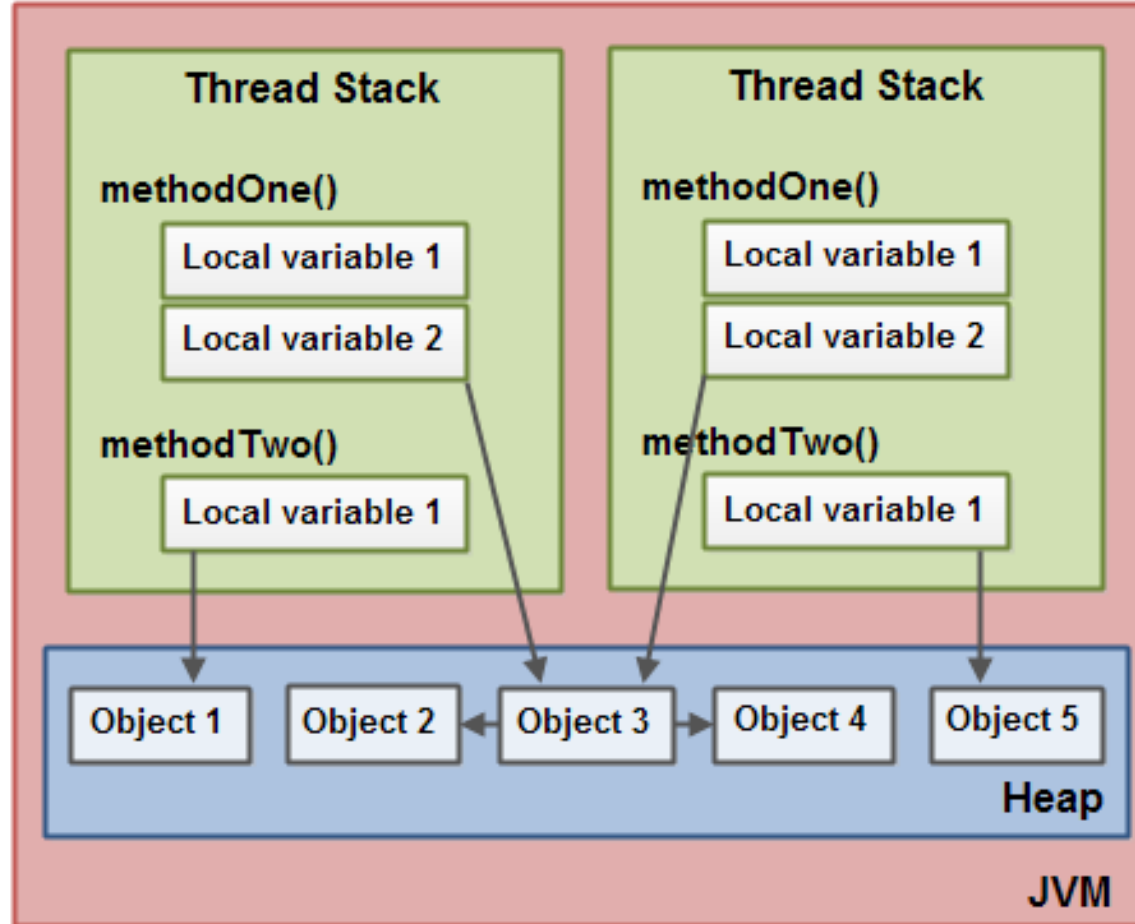


Unit of Concurrency

- **Multi-processing** - Multiple Processors/CPU's executing concurrently (Unit: CPU)
- **Multi-tasking** - Multiple tasks/processes running concurrently on a single CPU.
 - OS executes these tasks by switching between them very frequently (Unit: Process)
- **Multi-threading** - Multiple parts of the same program running concurrently.
 - Dividing the same program into multiple parts/threads and run those threads concurrently (Unit: Thread)

Processes and Threads

- A **process** is a program in execution running independently and isolated from others
- A **thread** is a path of execution within a process
 - It has its own call stack but can access shared data of other threads in the same process.
- A Java application runs by default in one process
- Within a Java application might have several threads to achieve parallel processing or asynchronous behavior



Improvements with concurrency



- Concurrency promises to perform certain tasks faster
 - A task = Several subtasks,
 - These subtasks can be executed in parallel
 - Save time (Better CPU Utilization)
- The theoretical possible performance gain can be calculated by *Amdahl's Law*

Amdahl's Law Illustrated

- If **B** is the percentage of the program which can not run in parallel and **N** is the number of processes, then the maximum performance gain is $1 / (B + ((1-B)/N))$.



B = Non-parallelizable
1 - B = Parallelizable



Saving time

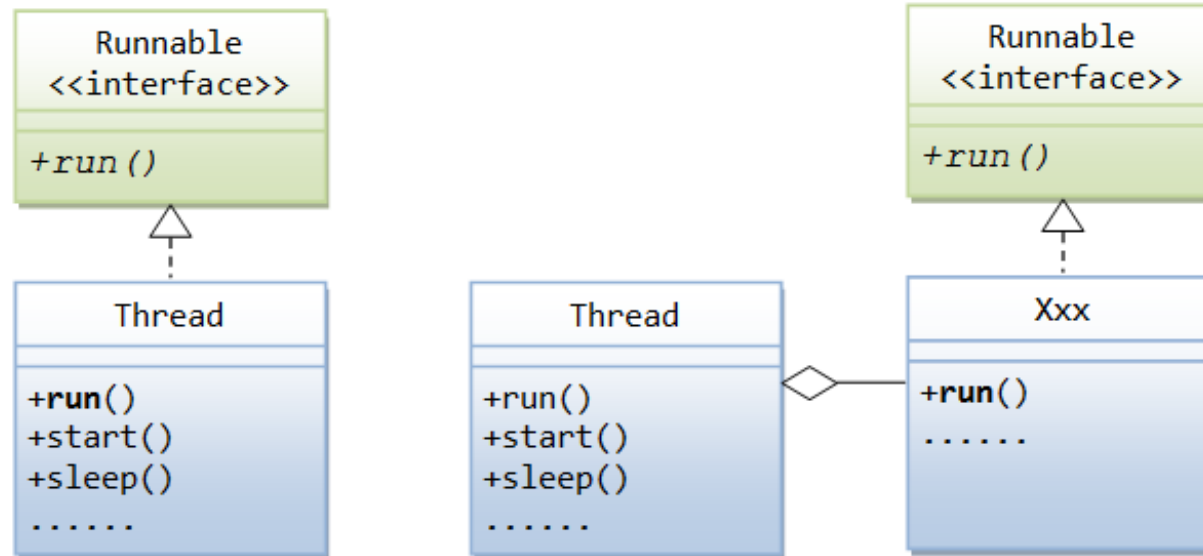
```
5 seconds reading file A
2 seconds processing file A
5 seconds reading file B
2 seconds processing file B
-----
14 seconds total
```

```
5 seconds reading file A
5 seconds reading file B + 2 seconds processing file A
2 seconds processing file B
-----
12 seconds total
```

Java Thread

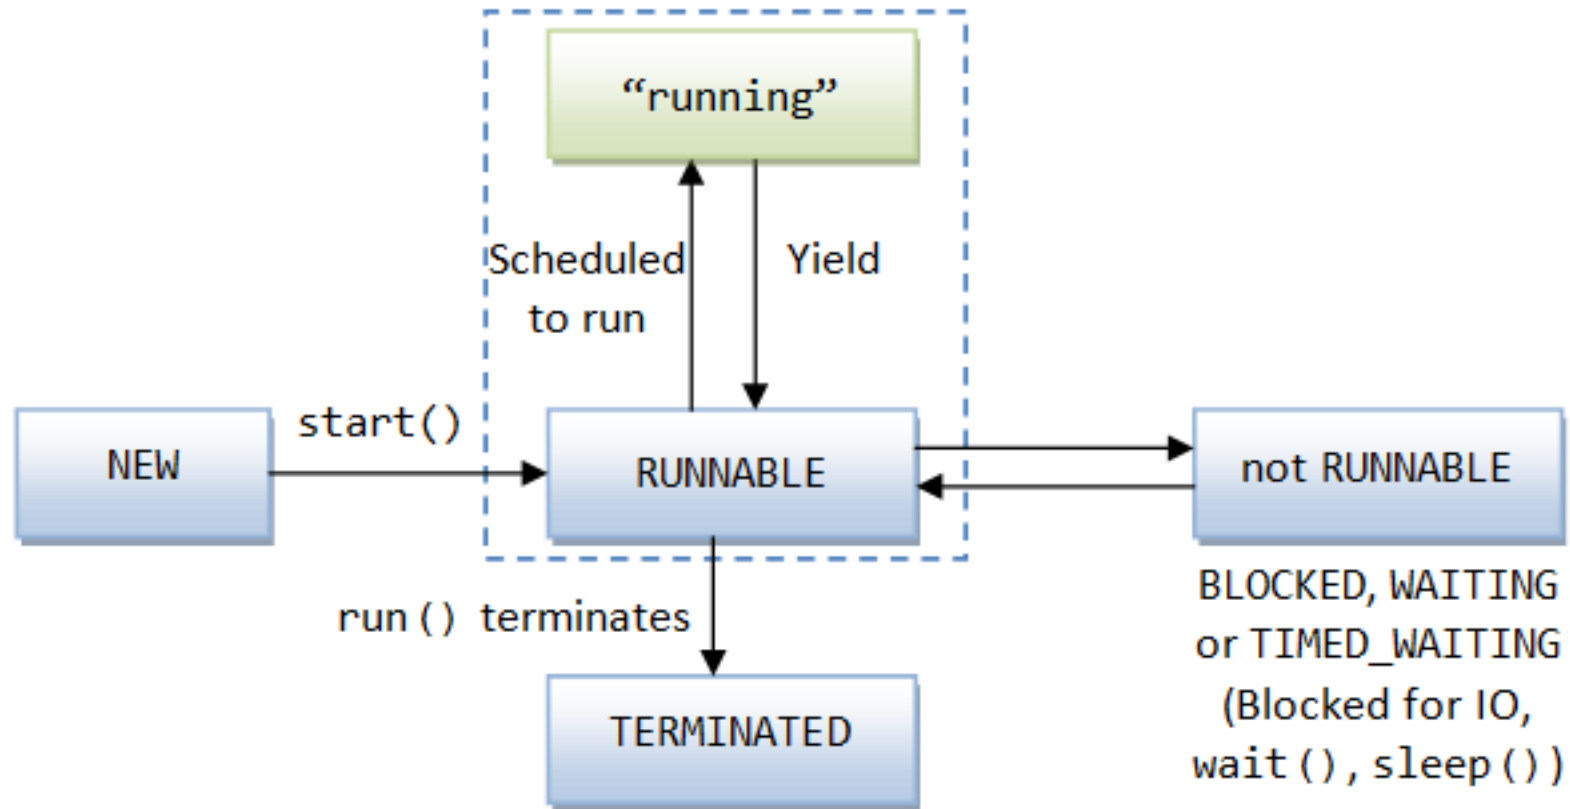
Creating and Starting a Thread

- By extending Thread class
- By providing a Runnable object
 - Runnable Anonymous Class
 - Runnable Lambda Expression




Runnable or Thread, which
one to use?

The Life Cycle of a Thread



Executor Framework

Executor Framework



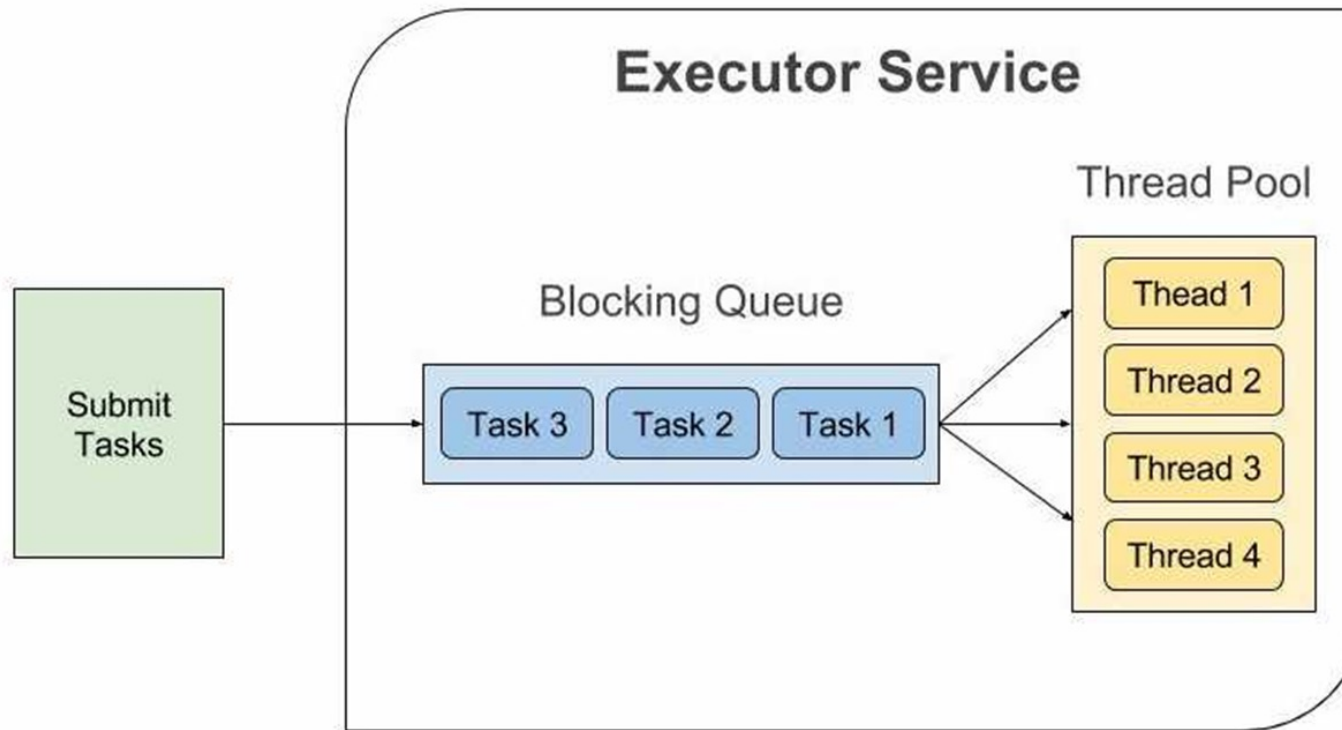
Executor Framework - Functionalities

- 1. Thread Creation:** Methods for creating threads, a pool of threads, the application can use to run tasks concurrently.
- 2. Thread Management:** Managing the life cycle of the threads in the thread pool.
 - You don't need to worry about whether the threads in the thread pool are active or busy or dead before submitting a task for execution
- 3. Task submission and execution:** Methods for submitting tasks for execution in the thread pool and deciding when the tasks will be executed

Java executor interfaces

- **Executor** - A simple interface that contains a method called `execute()` to launch a task specified by a `Runnable` object.
- **ExecutorService** - A sub-interface of `Executor` that adds functionality to manage the lifecycle of the tasks. It also provides a `submit()` method whose overloaded versions can accept a `Runnable` and a `Callable` (discussed later).
- **ScheduledExecutorService** - A sub-interface of `ExecutorService`. It adds functionality to schedule the execution of the tasks.

Thread Pool



Callable

Callable

- Runnable object to define the tasks that are executed inside a thread
- **What if you want to return a result from your tasks?**
- Java provides a Callable interface
- A Callable is similar to Runnable, except it can return a result and throw a checked exception
- Callable interface has a single method `call()` to contain the code executed by a thread

Callable examples

```
Callable<String> callable = new Callable<String>() {  
    @Override  
    public String call() throws Exception {  
        // Perform some computation  
        Thread.sleep(2000);  
        return "Return some result";  
    }  
};
```

```
//a lambda expression  
Callable<String> callable = () -> {  
    // Perform some computation  
    Thread.sleep(2000);  
    return "Return some result";  
};
```


Executing Callable tasks

- A Callable can be submitted to an executor service for execution
 - What about the Callable's result?
 - How do you access it?
 - When the result of the submitted task will be available
- Future can be used to fetch the result of the task when it is available

Cancelling a Future

- Cancelling a future by using `Future.cancel()` method
- The `cancel()` method accepts `mayInterruptIfRunning`.
 - `True`: the thread that is currently executing the task will be interrupted
 - `False`: in-progress tasks will be allowed to complete.
- `isCancelled()`: check if a task is canceled or not.
- After the cancellation of the task, `isDone()` will always be true.

invokeAll & invokeAny

Submit multiple tasks and wait for all of them to complete

- Executing multiple tasks by passing a collection of Callables to the `invokeAll()` method
 - `invokeAll()` returns a list of Futures
 - Any call to `future.get()` will be blocked until all the Futures are complete

Submit multiple tasks and wait for any one of them to complete

- The `invokeAny()` method accepts a collection of Callables and returns the result of the fastest Callable.

Synchronization

Issues with concurrency

- Threads have their call stack but can also access shared data
- **Access problem:** if several threads access and change the same shared data at the same time
 - Safety failure (inconsistent data)
 - Thread Interference Errors (Race conditions)
- **Visibility problem:** if thread A reads shared data, which is later changed by thread B and thread A is unaware of this change
 - Liveness failure (e.g., deadlocks)
 - Memory Consistency Errors

Synchronization

How do we avoid those problems?

1. Only one thread can read and write a shared variable at a time.
 - When one thread is accessing a shared variable, other threads should wait until the first thread is done.
 - This guarantees that the access to a shared variable is `Atomic`, and that multiple threads do not interfere.
2. Whenever any thread modifies a shared variable, it automatically establishes a *happens-before* relationship with subsequent reads and writes of the shared variable by other threads.
 - This guarantees that changes done by one thread are visible to others.

Synchronized Methods

- The `synchronized` keyword makes sure that only one thread can enter the sync methods at one time

Synchronized Blocks

- Java internally uses a so-called *intrinsic lock* or *monitor lock* to manage thread synchronization. Every object has an intrinsic lock associated with it.

Volatile Keyword

- `Volatile` keyword is used to avoid memory consistency errors in multithreaded programs.
 - It tells the compiler to avoid doing any optimizations to the variable.
 - If you mark a variable as `volatile`, the compiler won't optimize or reorder instructions around that variable.

Locks

- ReentrantLock is a mutually exclusive lock with the same behavior as the intrinsic/implicit lock accessed via the synchronized keyword.
 - Thread that currently owns the lock can acquire it more than once without any problem.
- The ReentrantLock also provides various methods for more fine-grained control
- The tryLock() method tries to acquire the lock without pausing the thread.
 - If the thread couldn't acquire the lock because it was held by some other thread, then it returns immediately instead of waiting for the lock to be released.

ReadWriteLock

- ReadWriteLock consists of a pair of locks - one for read access and one for write access.
 - The read lock may be held by multiple threads simultaneously as long as the write lock is not held by any thread.
- ReadWriteLock allows for an increased level of concurrency.
 - It performs better compared to other locks in applications where there are fewer writes than reads.

Atomic Variables

- Java's concurrency API defines several classes in `java.util.concurrent.atomic` package that support Atomic operations on single variables.
- Atomic classes internally use compare-and-swap instructions supported by modern CPUs to achieve synchronization. These instructions are generally much faster than locks.
- The `AtomicInteger.incrementAndGet()` method is atomic,
 - Safely call it from several threads simultaneously and ensure the access to the count variable will be synchronized.

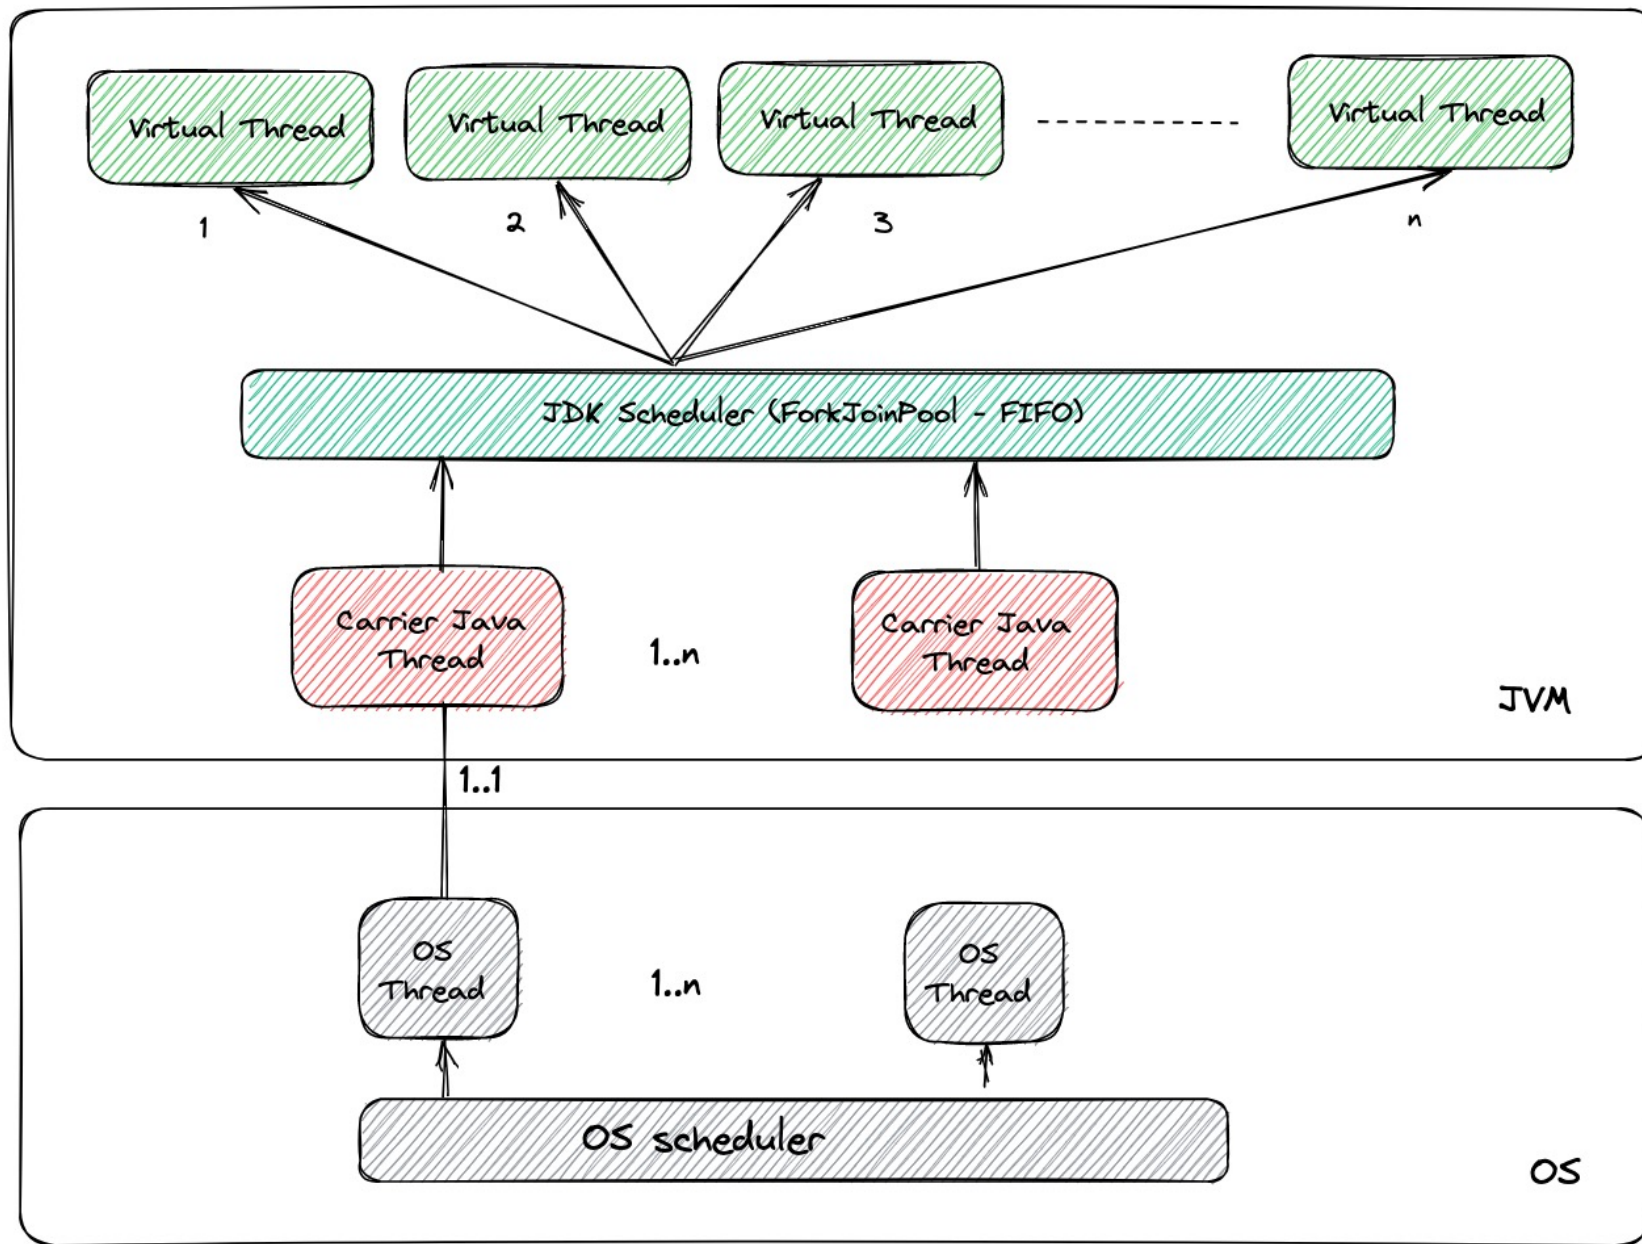
Java Virtual Thread

Java (normal) Thread

- To create a new kernel thread, we must do a system call, and that's a costly operation
 - Using thread pools instead of reallocating and deallocating threads as needed
- Scaling up application by adding more threads
 - Context switching + Memory footprint
 - Cost of thread maintaining may be significant and affect the processing time

Java Virtual Thread

- Virtual threads are managed by the JVM
- Their allocation doesn't require a system call
- They're free of OS's context switch
 - Virtual threads run on the *carrier kernel thread* used under-the-hood
 - More virtual threads...
 - Managing virtual threads is much cheaper



API is the same!!!

```
Runnable printThread = () -> System.out.println(Thread.currentThread());

ThreadFactory virtualThreadFactory = Thread.builder().virtual().factory();
ThreadFactory kernelThreadFactory = Thread.builder().factory();

Thread virtualThread = virtualThreadFactory.newThread(printThread);
Thread kernelThread = kernelThreadFactory.newThread(printThread);

virtualThread.start();
kernelThread.start();
```